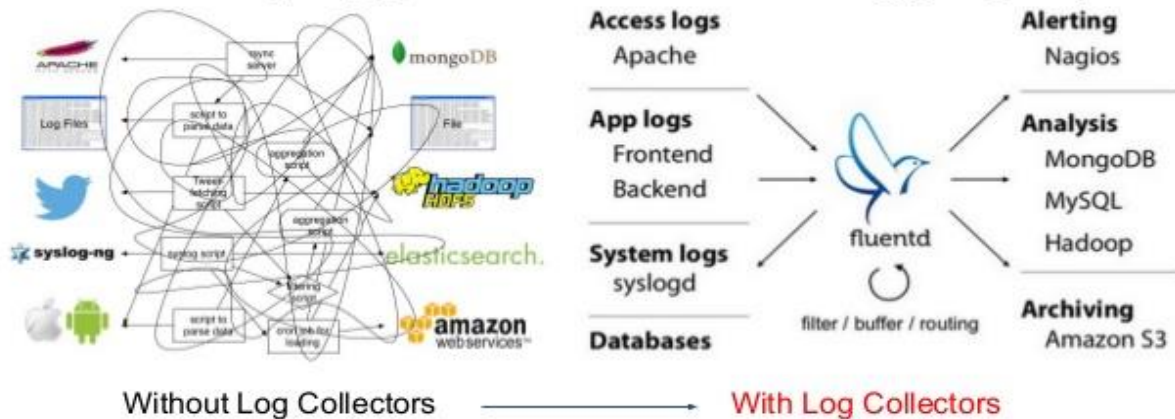


## What are Log Collectors?

- Provide pluggable and unified logging layer



Images from <http://fluentd.org/>

3

## Two Popular Log Collectors

- Fluentd
  - Written in CRuby
  - Used in Kubernetes
  - Maintained by Treasure Data Inc.
- Logstash
  - Written in JRuby
  - Maintained by elastic.co
- They have similar features
- Which one is better for you?



5



# fluentd

Structured logging

Reliable forwarding

Pluggable architecture

<http://fluentd.org/>

## WHAT'S FLUENTD?

Buffering, HA (failover),  
load balancing, etc.

**AN EXTENSIBLE & RELIABLE DATA COLLECTION  
TOOL**

Simple core + plugins

Like syslogd

## **Reliability (core + plugin)**

- > Buffering**
  - > Use file buffer for persistent data**
  - > buffer chunk has ID for idempotent**
- > Retrying**
- > Error handling**
  - > transaction, failover, etc on forward plugin**
  - > secondary**

## CORE

- Divide & Conquer
- Buffering & Retries
- Error Handling
- Message Routing
- Parallelism

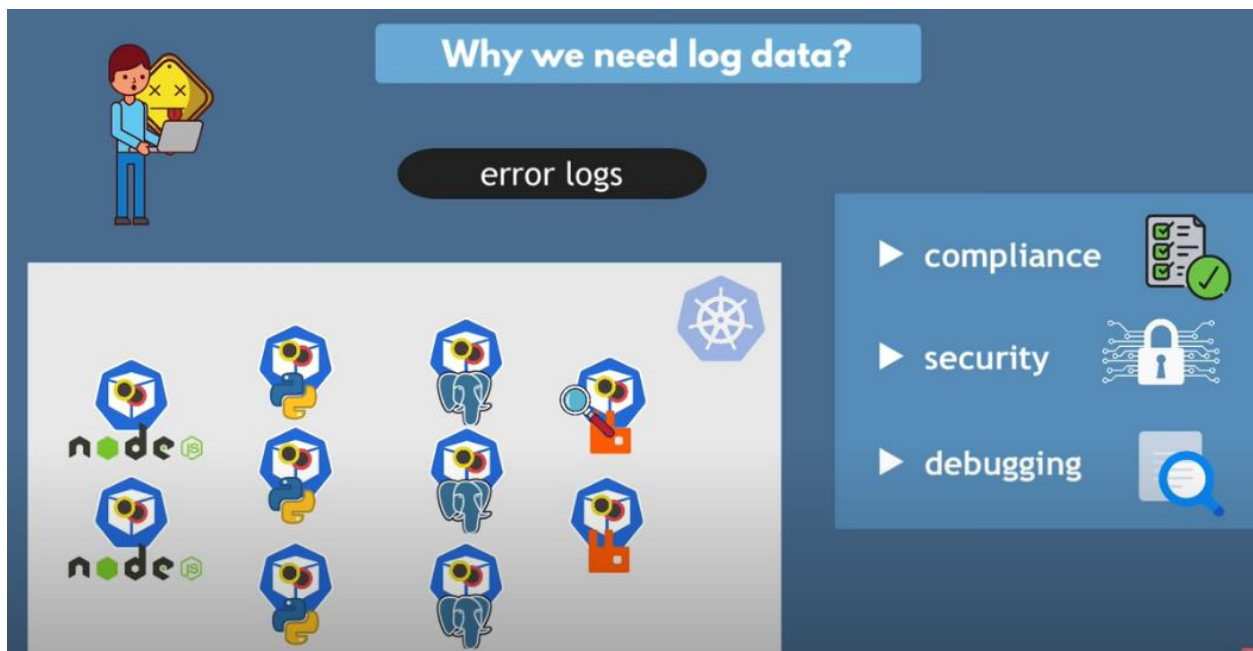
Common concerns

## PLUGINS

- Read Data
- Parse Data
- Buffer Data
- Write Data
- Format Data

Use case specific

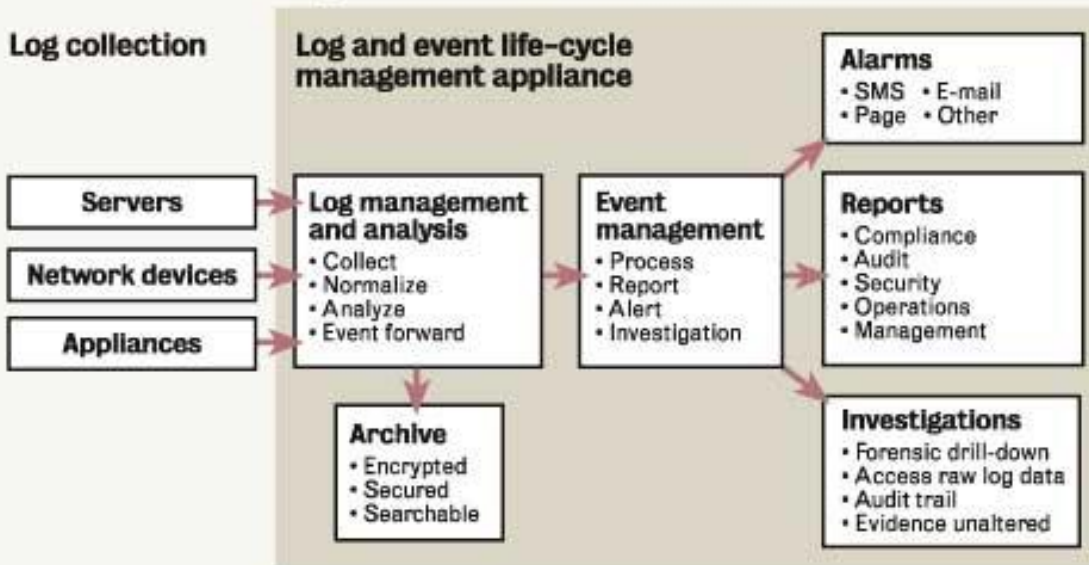
Why we actually need log data?



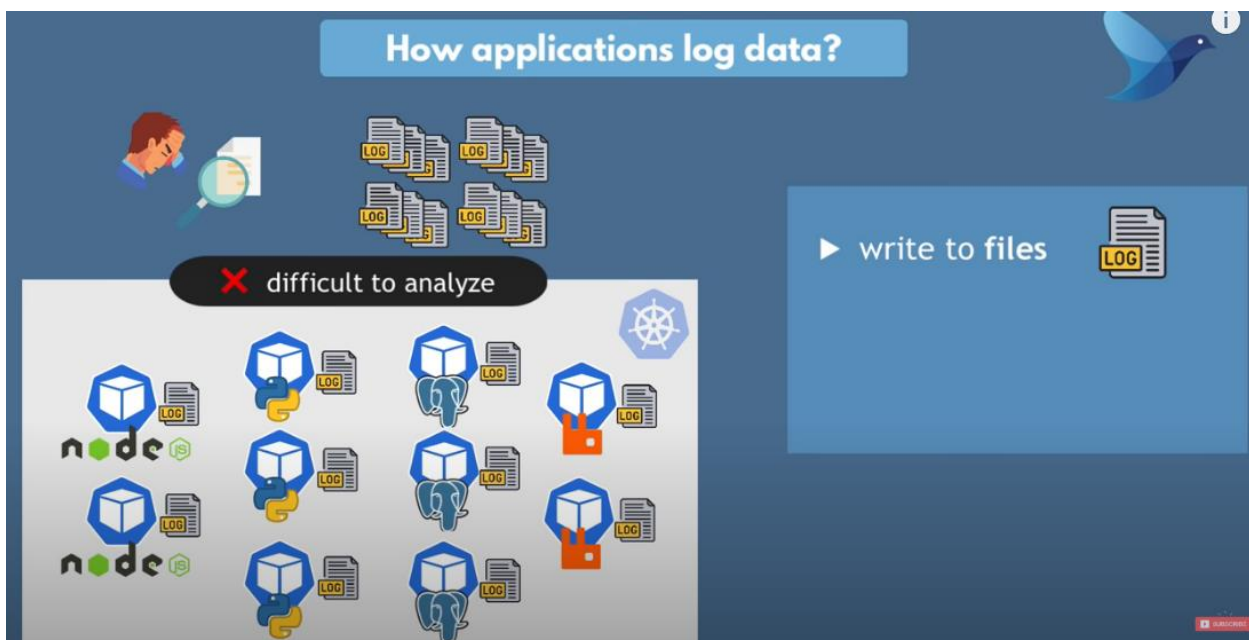
## How it works

A log and event management appliance should be able to collect from multiple sources, normalize, time stamp and analyze the results, generate alerts and reports for audit and regulatory compliance, and support data mining for forensic and root-cause investigations.

### Log collection



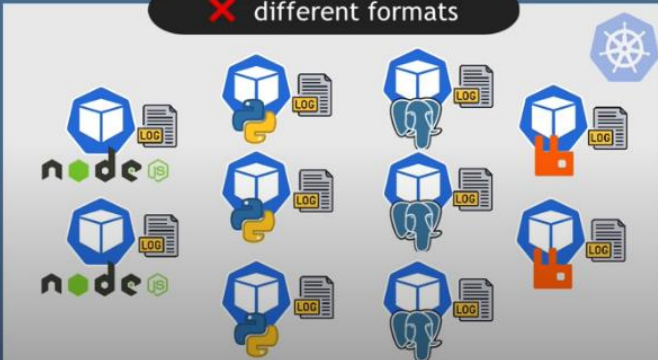
## How applications log data?



## How applications log data?



✗ different formats



► write to files



## How applications log data?



✗ each app must be configured



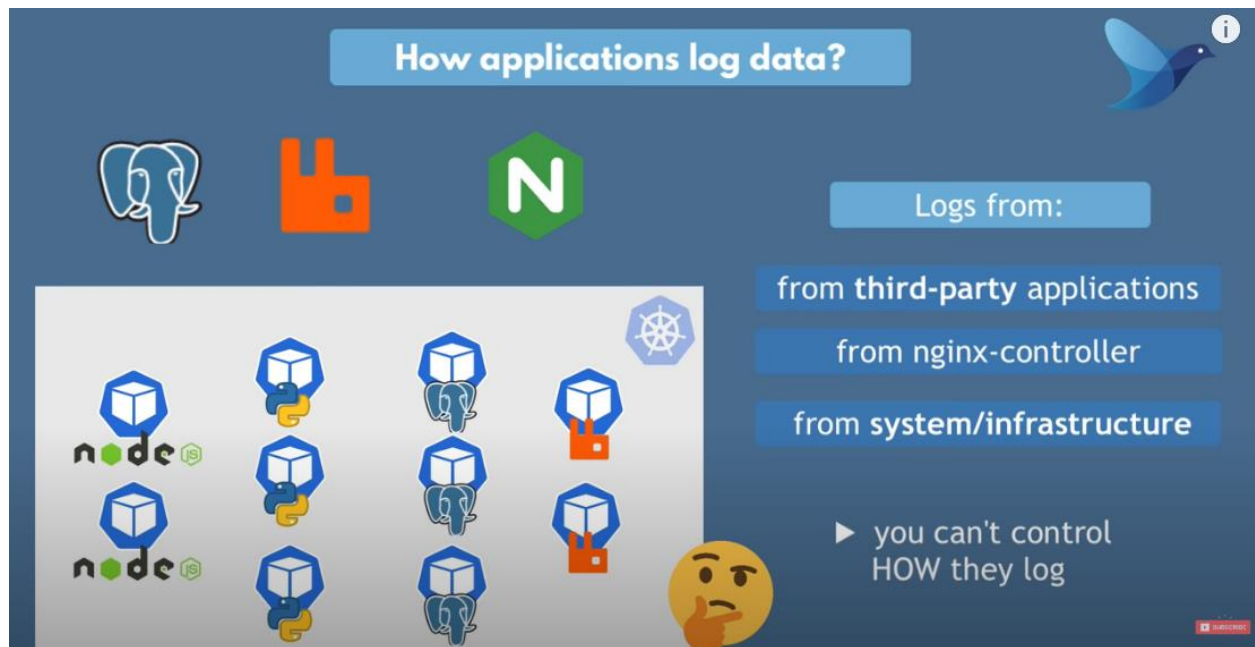
► write to files



► log directly into DB







### problem 1: need ad-hoc parsing

the text-based logs have their own format, and analytics engineer need to write a dedicated parser for each format. but that's probably not the best use of your time. you should be analyzing data to make better business decisions instead of writing one parser after another.

### problem 2: lacks freshness

the logs lag, the realtime analysis of user behavior makes feature iterations a lot faster. a nimbler a/b testing will help you differentiate your service from competitors.

this is where **fluentd** comes in. **we believe fluentd solves all issues of scalable log collection** by getting rid of files and turning logs into true semi-structured data streams.

## what's fluentd?

the best way to describe fluentd is 'syslogd that understands json'. the notable features are:

- **easy installation** by rpm/deb/gem
- **small footprint** with 3000 lines of ruby
- **semi-structured data** logging
- **easy start** with small configuration
- **fully pluggable** architecture, and plugin distribution by ruby gems

# Fluentd


- Fluentd is an open source data collector for unified logging layer.
- Fluentd allows you to unify data collection and consumption for a better use and understanding of data.

Fluentd is an open-source log data collection software designed to separate data sources from the back-end system using a Unified Logging Layer (ULL). Here are some of the software's major highlights:

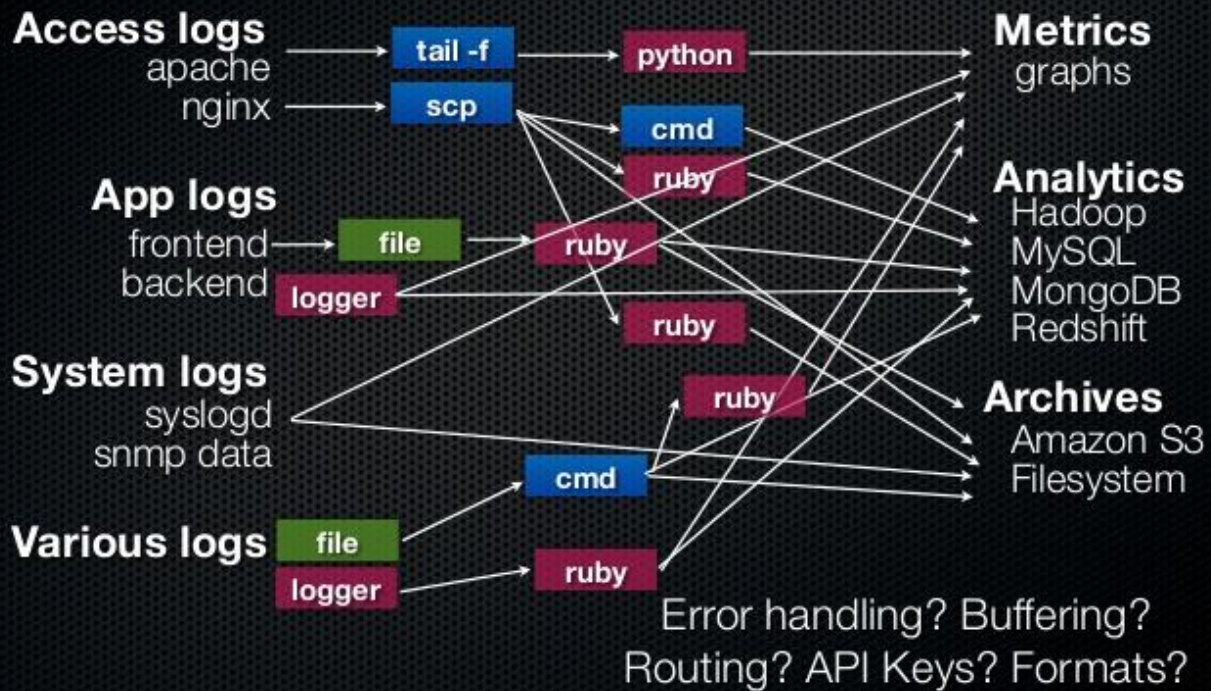
- **Unified Log Structuring With JSON:** Fluentd structures your log data in JSON, allowing you to collect, filter, buffer, and centralize logs from multiple sources (mobile devices, web servers, etc.).
- **Flexible Plug-In Architecture:** Using Fluentd plug-ins, IT teams can expand their functionalities and easily connect with several data sources and outputs.
- **Minimal Resource Requirements:** Fluentd is a lightweight log collection tool. It's written using C language and a Ruby framework, which is why it requires minimal system resources and memory allocation.
- **Data Loss Prevention:** Fluentd supports robust failover to avoid data loss. It's an extremely reliable tool in terms of memory allocation and file-based buffering, and it can store data in multiple systems.

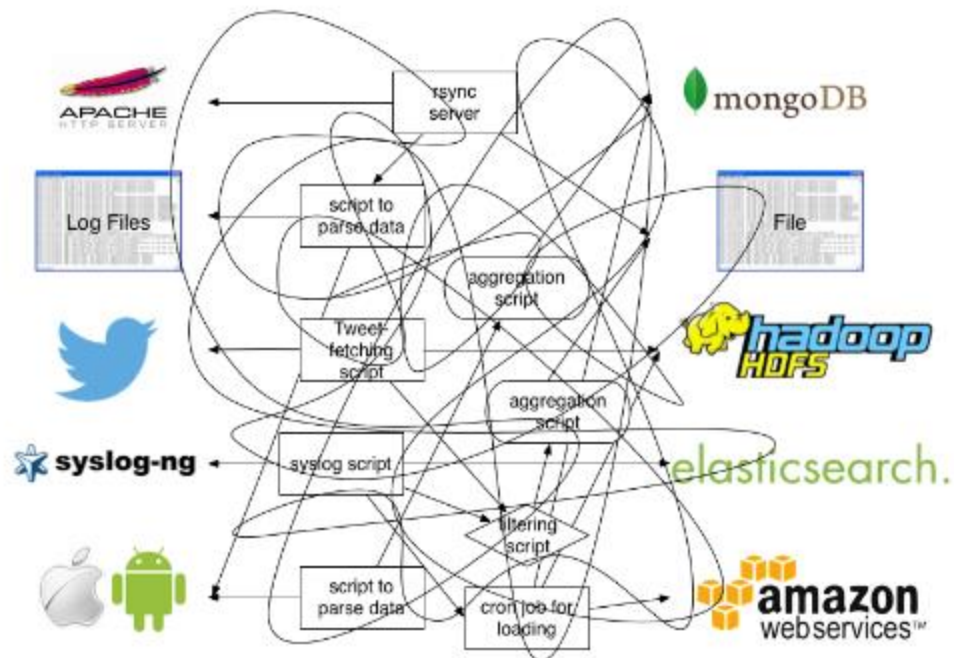


## Before Fluentd



# Before Fluentd: **CHAOS**






Before Fluentd

Before EFK

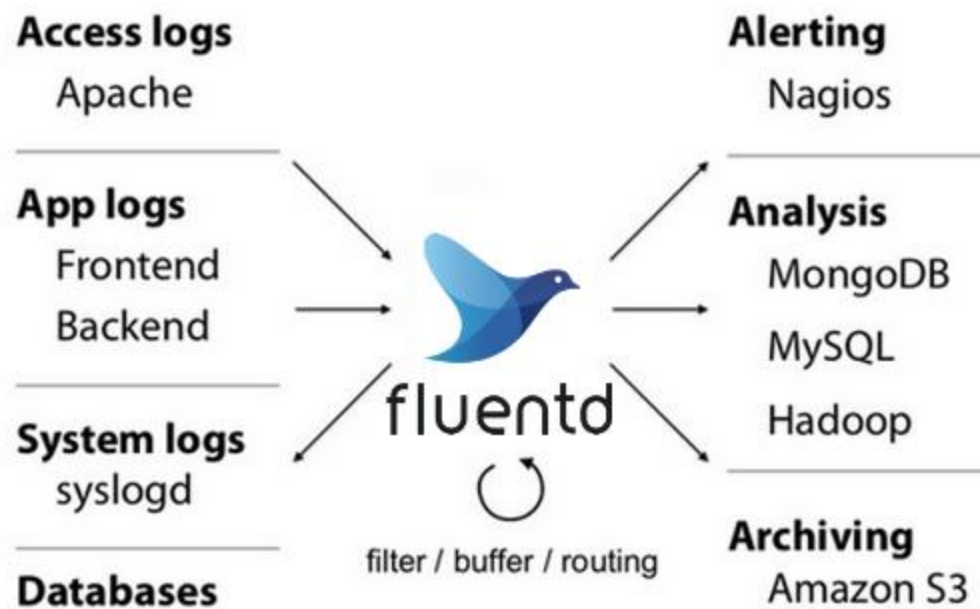
- Log files across the server
- Grep/shell scripting mess
- Production system access

## After Fluentd



EFK:

- Detecting log patterns
- Easily Searchable
- Trouble-shooting issues



After Fluentd



# After Fluentd: **Controllable**

## **Access logs**

apache  
nginx

## **App logs**

frontend  
backend

## **System logs**

syslogd  
snmp data

## **Various logs**

## **Metrics**

graphs

## **Analytics**

Hadoop  
MySQL  
MongoDB  
Redshift

## **Archives**

Amazon S3  
Filesystem

Fluentd does:

**Format, Buffer, Retry, Route**

## Log Management Challenges



Ever increasing Volume,  
Variety, Velocity of Logs



Logs from many disparate  
sources



Unable to extract  
actionable insights from  
logs



Unable to detect security  
threats easily.  
Compliance/Audit Risks



Stringent short-term and  
long-term log archival  
requirements



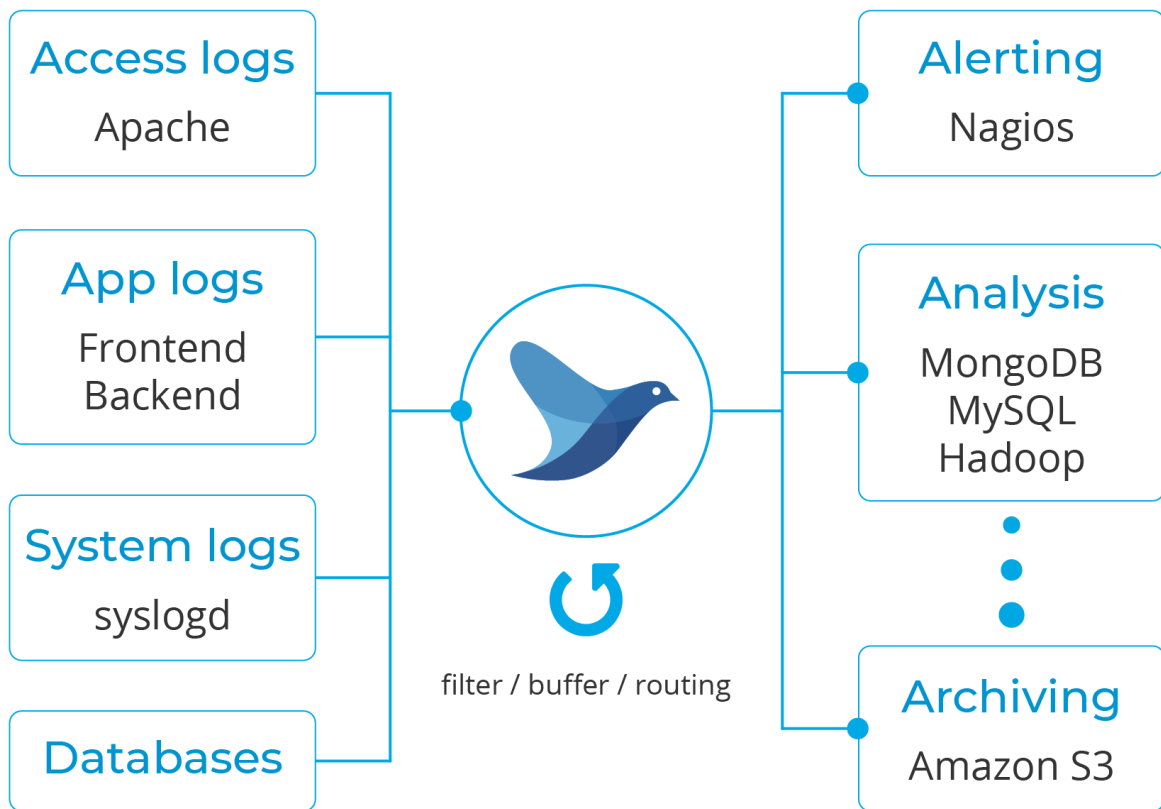
## Need for Log Management

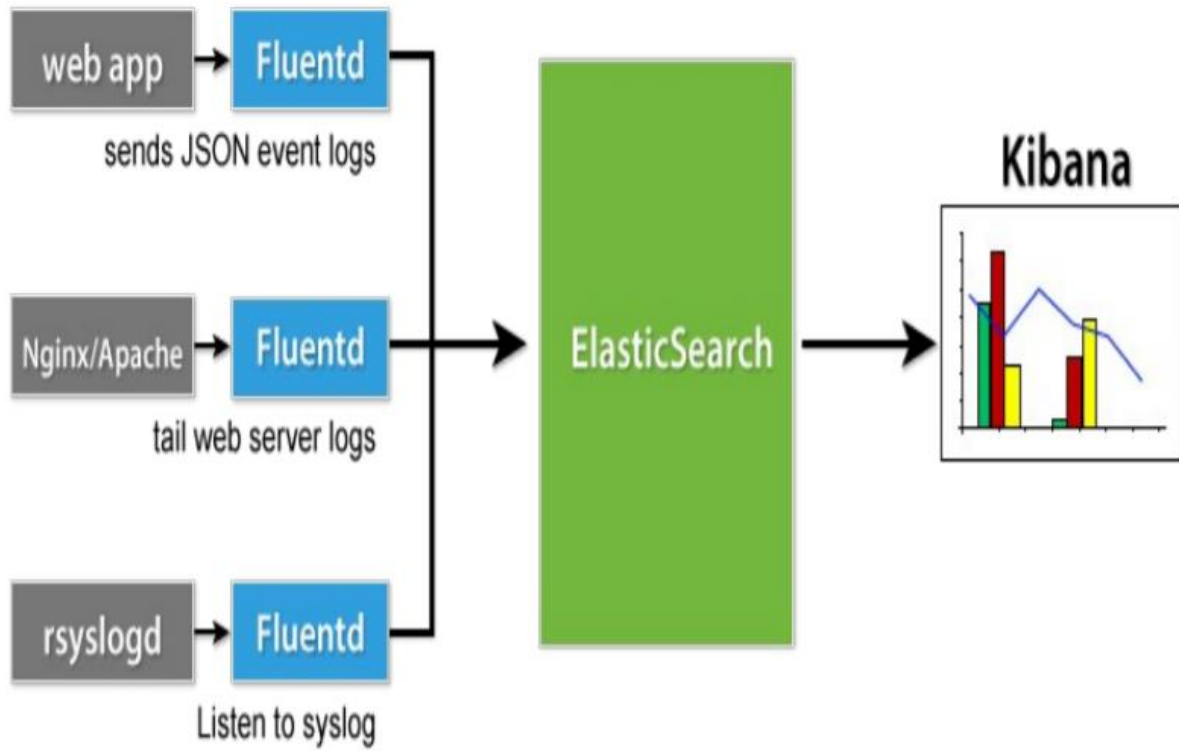
- Logs are usually in proprietary format and difficult to manage
- Routine log reviews and analysis are beneficial for identifying security incidents, policy violations, fraudulent activity, and operational problems
- Logs can also be useful for performing auditing and forensic analysis, supporting the organization's internal investigations, establishing baselines, and identifying operational trends
- Legal compliance. For critical applications like, health, public financial records, bank accounts, Government requires the organizations to maintain logs
- Protecting the trustworthiness of the log sources and also, the logs themselves need to be protected from malicious activities

# Log Management Infrastructure

---


- A three-tier Architecture
  - Log generation : Synchronized hosts generate
  - Logs analysis and storage : One or more log servers that receive the logged data. This transfer is either real-time or periodic. Such servers are called collectors or aggregators
  - Log monitoring : analyze and monitor the logged data using application consoles



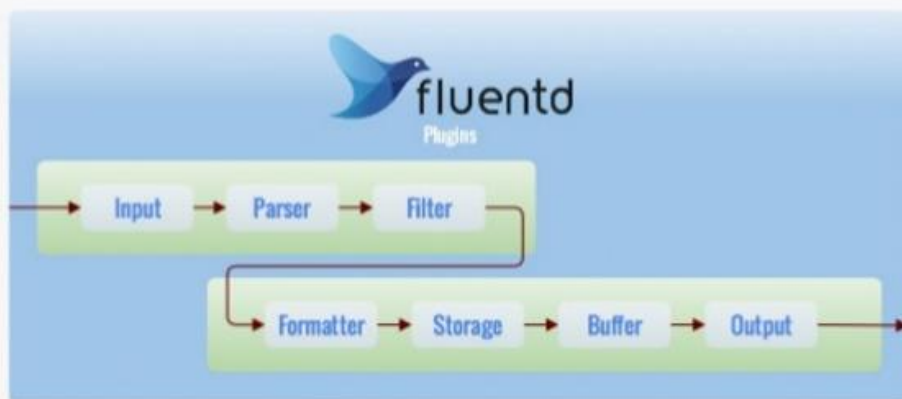





## Architecture



## Available Fluentd Plugin Types



## INTERNAL ARCHITECTURE



### input

input is the place where the log comes in. the user can extend it to feed the events from various sources. the example input supported officially includes: http+json, tailing files (apache log parser is supported), syslog. of course you can add input plugin by writing a ruby plugin.

### buffer

buffer exists for reliability. when the output fails, the events are kept by buffer and automatically retried. memory or file buffer is supported now.

### output

buffer creates chunks of logs, and passes them to the output. output stores or forwards chunks. the buffer waits several seconds to 1 minute, to create chunks. this is really efficient for writing into the storage which supports batch-style importing.

## Fluentd Plugins Used

- `in_forward`: capture logs securely on port 24224 and unsecurely on port 24223
- `parser_multi_format`: parse logs where the log stream has more than one format e.g Redis
- `filter_record_transformer`: used to add a 'source' key value pair
- `out_elasticsearch`: forward logs to Elasticsearch targetting different indices as appropriate
- `out_copy`: copies logs to more than one output source e.g. Elasticsearch AND stdout
- `out_rewrite_tag_filter`: used to rewrite the tags from k8s and re-emit logs to process

## Fluentd

- Fluentd has 5 types of plugins: Input, Parser, Output, Formatter and Buffer

### Input Plugins:

- `in_forward`
- `in_http`
- `in_tail`
- `in_exec`
- `in_syslog`

## Output Plugins:

- out\_forward
- out\_mongo
- out\_file
- out\_s3


## Buffer Plugins:

Buffer plugins are used by buffered output plugins, such as out\_file, out\_forward, etc

- buf\_memory
- buf\_file

## Configuration: Fluentd

- Every inputs are tagged
- Logs will be routed by tag



## Fluentd Configuration: Input

- Every inputs will be tagged

```
<source>  
  @type tail  
  path /var/log/nova/nova-api.log  
  tag openstack.nova  
</source>
```

Example of tailing nova-api log

## Fluentd Configuration: Output

```
<match openstack.nova> # nova related logs  
  @type elasticsearch  
  host example.com  
</match>  
  
<match openstack.*> # all other OpenStack related logs  
  @type influxdb  
  # ...  
</match>
```

Routed by tag  
(First match is priority)

Wildcards can be used



# Fluentd Configuration: Copy

```
<match openstack.*>
```

```
  @type copy
```

```
  <store>
```

```
    @type influxdb    tag: openstack.*
```

```
  </store>
```


```
  <store>
```

```
    @type elasticsearch
```

```
  </store>
```

```
</match>
```


Copy plugin enables multiple outputs for a tag



10

# Fluentd Transport: forward

- Active-Active (Load Balancing)




```
<match openstack.*>
  type forward
  <server>
    host dest1
  </server>
  <server>
    host dest2
  </server>
</match>
```

23

# Fluentd Transport: forward

- Active-Standby




```
<match openstack.*>
  type forward
  <server>
    host primary
  </server>
  <server>
    host secondary
    standby
  </server>
</match>
```

24

# Fluentd Transport: forward

- Weighted Load Balancing




```
<match openstack.*>
  type forward
  <server>
    host dest1
    weight 60
  </server>
  <server>
    host dest2
    weight 40
  </server>
</match>
```

25

# Fluentd Transport: forward


- At-least-one Semantics  
(may affect performance)



Logs are re-transmitted  
until ACK is received

```
<match openstack.*>
  type forward
  require_ack_response
  <server>
    host dest
  </server>
</match>
```


**All Input events should be tagged:**





**It matches the tag against the outputs:**



**It can support multiple sources and destinations**

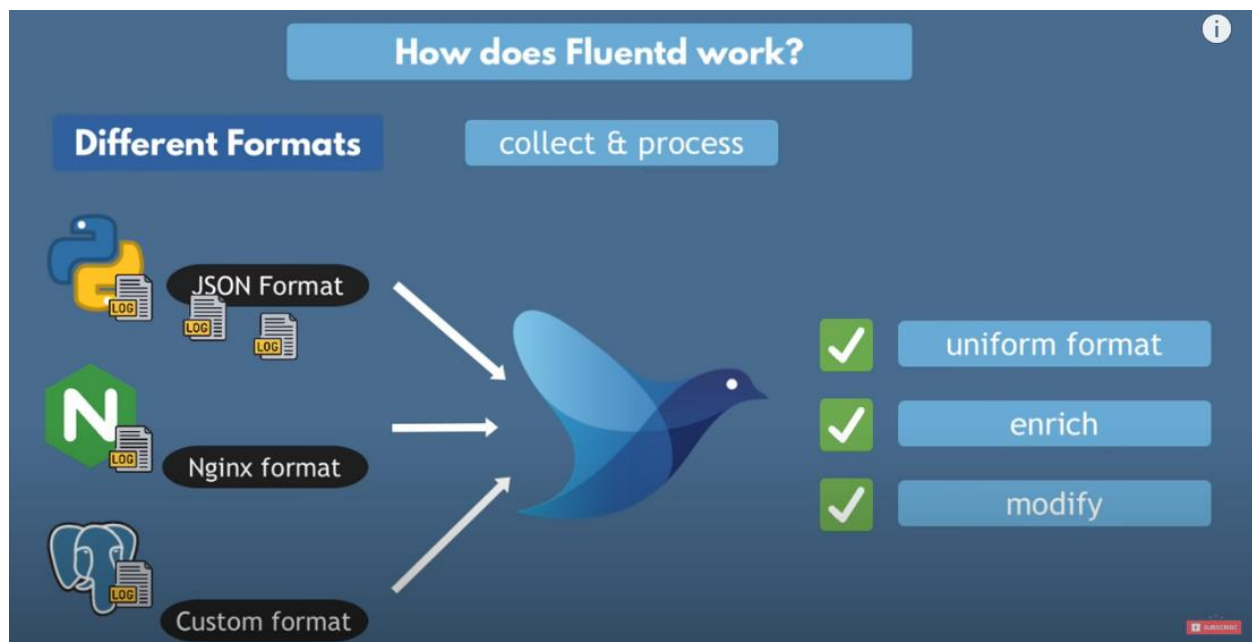


**Fluentd sends the event to matched output**

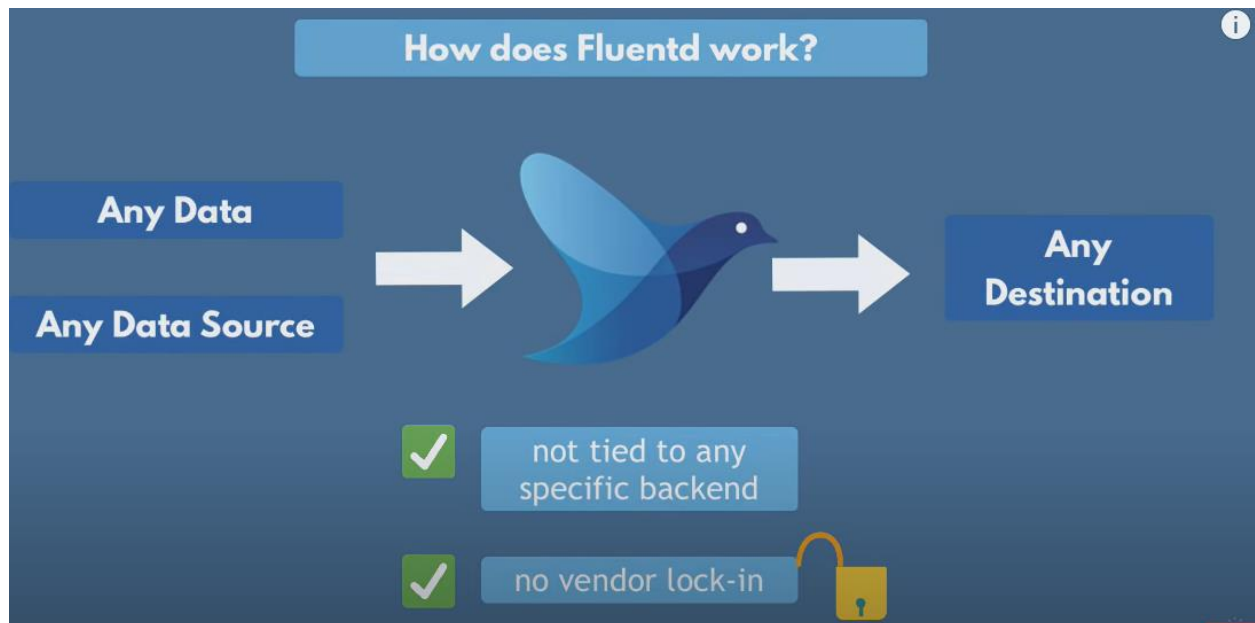
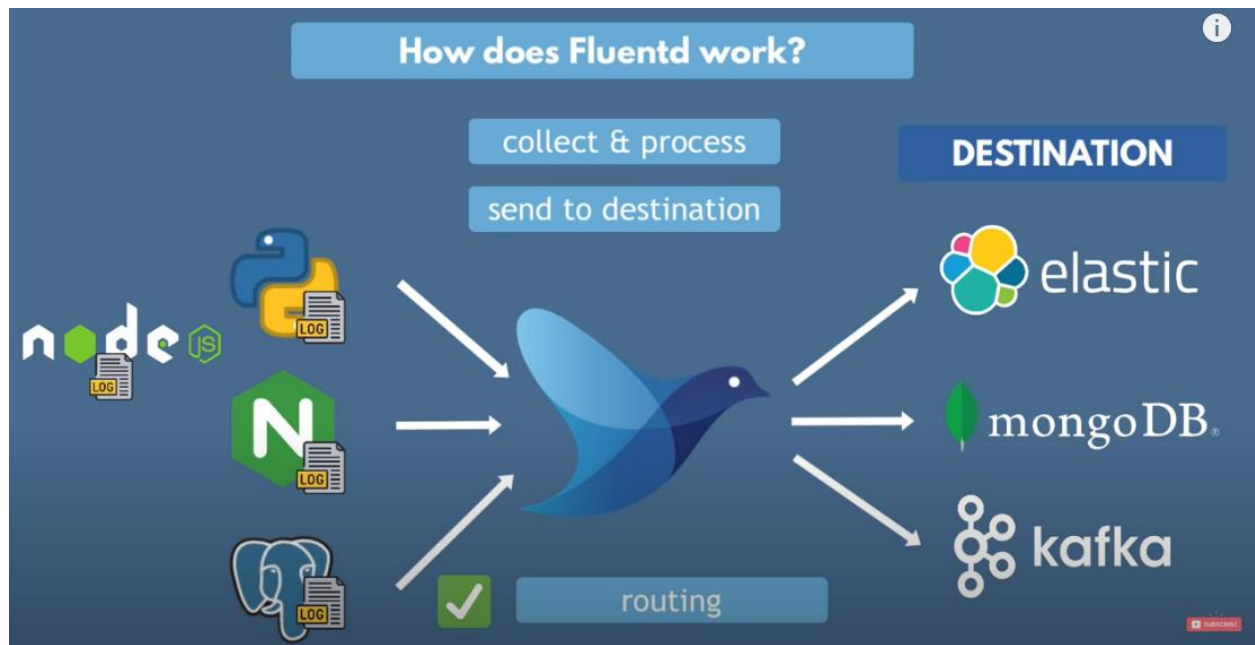


# Fluentd Event Structure

- **tag**: From where an event originated; used for message routing
- **time**: The epoch time at which an event occurred
- **record**: The event log content as a JSON object









## Built-In Reliability

- ▶ saves data on hard drive
- ▶ data is still there after restarts
- ▶ no additional storage configuration needed
- ▶ retries
- ▶ high availability through clustering

EFK( **E**lastic **F**luentD **K**ibana)



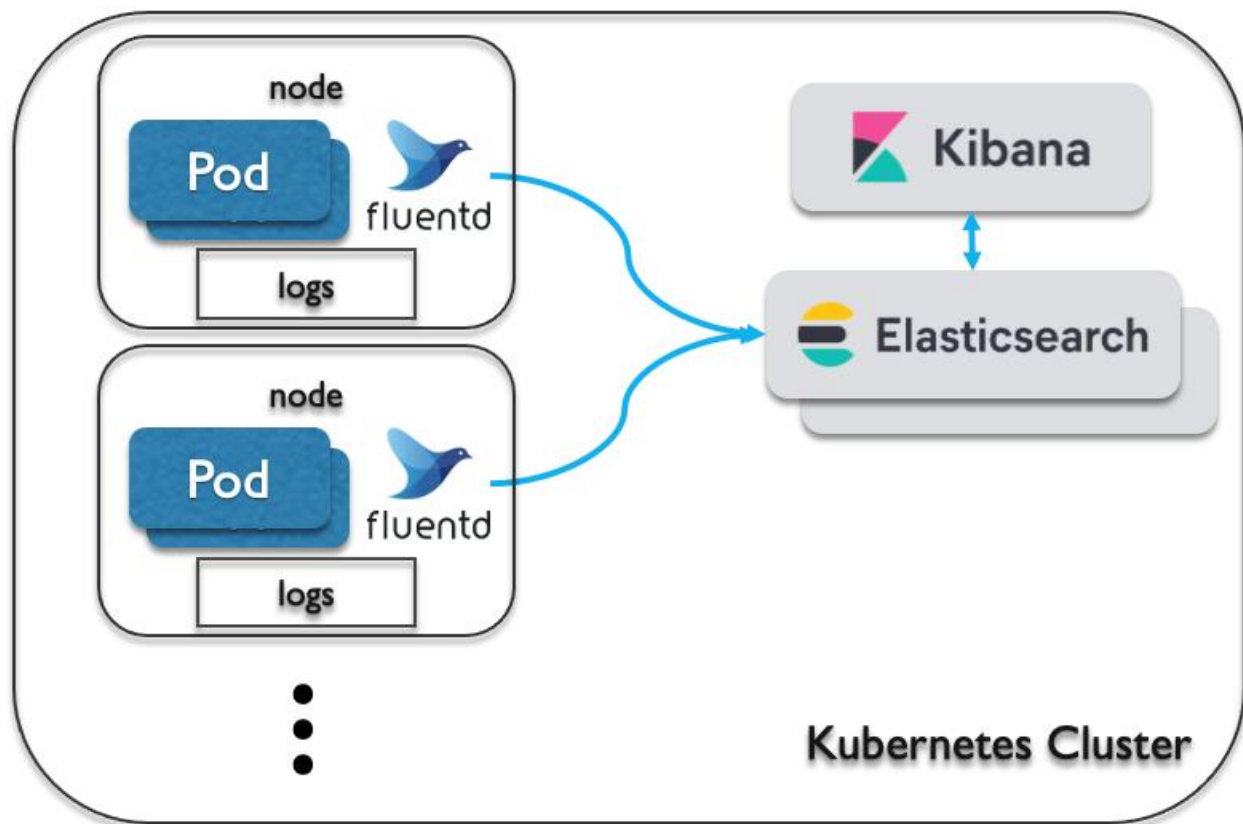
Logs



Elastic  
Search



Kibana

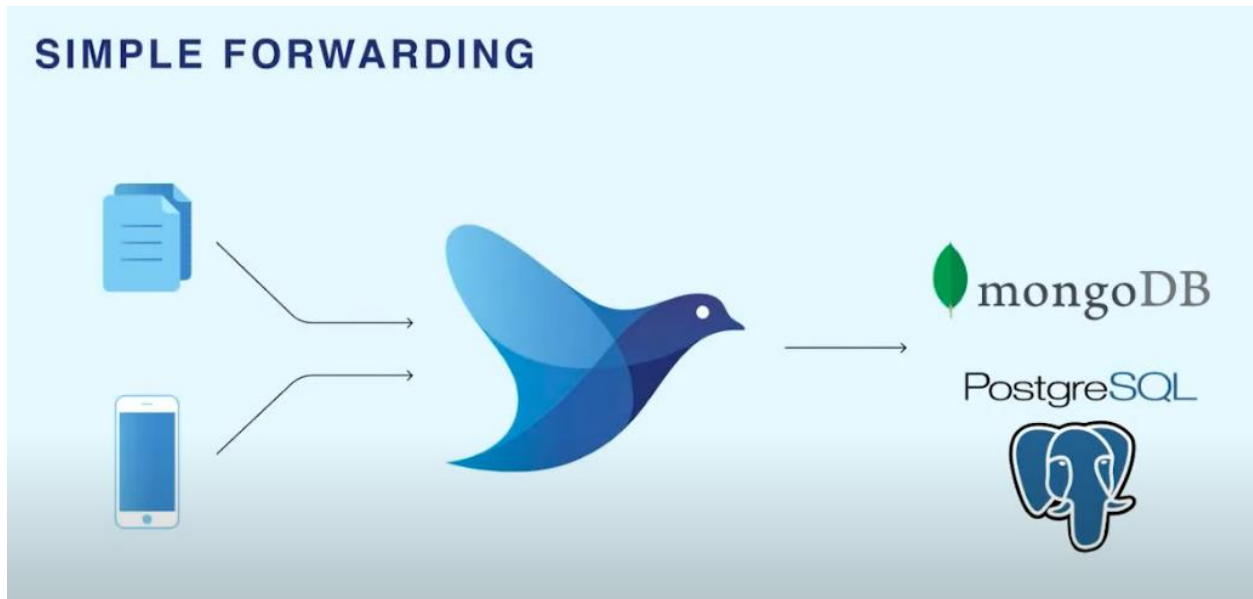


## Fluentd vs Logstash

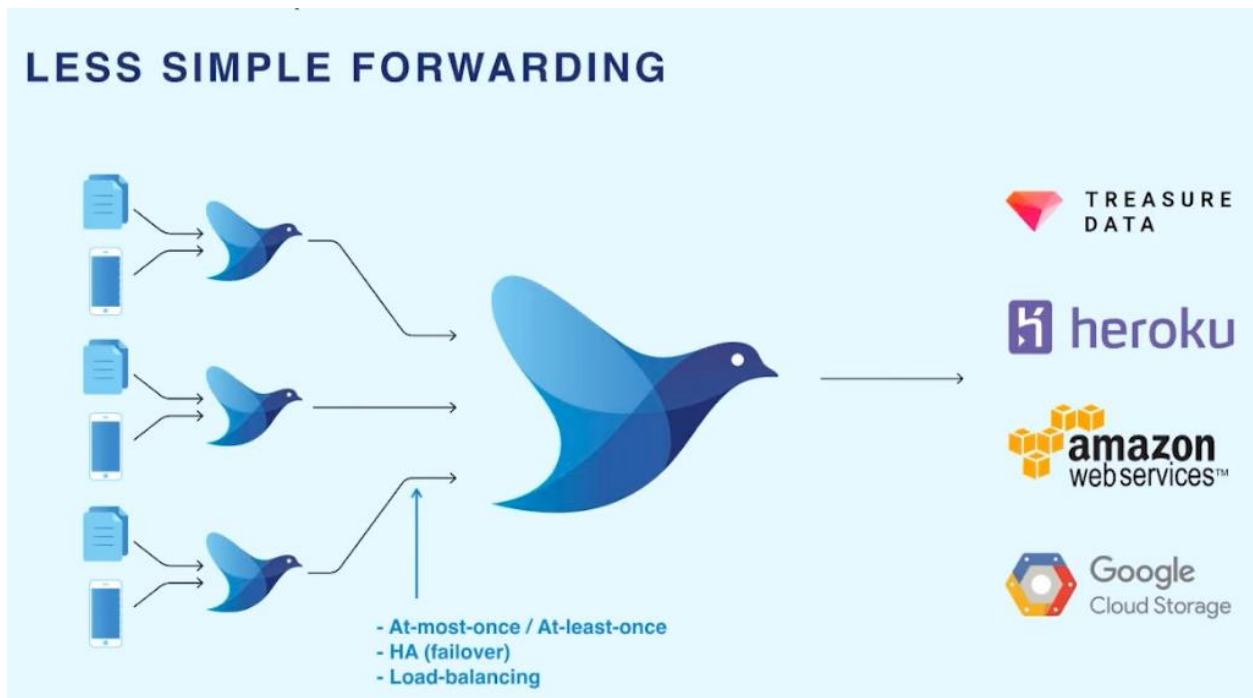
	Platform	Event Routing	Plugin Ecosystem	Transport	Performance
<b>Logstash</b>	Linux & Windows	Algorithmic statements	Centralized	Deploy with Redis for reliability.	Uses more memory. Use Elastic Beats for leafs.
<b>Fluentd</b>	Linux & Windows	Tags	Decentralized	Built-in reliability but hard to configure.	Uses less memory. Use Fluent Bit and Fluentd Forwarder for leafs.

## USE CASES


### SIMPLE FORWARDING




### LESS SIMPLE FORWARDING



## LAMBDA ARCHITECTURE



## LOGGING TO GCP



## CONTAINER LOGGING

